EECS 151/251A ASIC Lab 1: Getting around the Compute Environment

Prof. Borivoje Nikolic and Prof. Sophia Shao
TAs: Cem Yalcin, Rebekah Zhao, Ryan Kaveh, Vighnesh Iyer

Overview

The process of VLSI design is somewhat different than developing software, designing analog circuits, or even FPGA-based design. Instead of using a single unified graphical program (eg. Eclipse, Cadence Virtuoso, or Xilinx ISE), VLSI design is done using dozens of command line interface based tools on a Linux machine. These tools primarily use text files as their inputs and outputs, and include graphical interfaces mainly for visualization and not for design. Therefore, high familiarity with Linux, text manipulation, file manipulation, and scripting is required to successfully complete any of the labs this year.

The goal of this lab is to introduce the basic techniques needed to use the computer aided design (CAD) tools that are taught in this class. Mastering the topics in this lab will help you save hours of time in later labs, and make you a much more effective chip designer. While you go through this lab, focus on how these techniques will allow you to automate tasks and improve your efficiency. Chip design requires plenty of trial-and-error, so the key to your success will be performing trials and identifying errors quickly.

Administrative Info

This lab, like all labs, will be turned in electronically using Gradescope. Please upload a *pdf* document with the answers to the questions throughout the lab.

Getting an Instructional Account

You are required to get an EECS instructional account to login to the workstations in lab. This can be done by using the WebAcct here: `http://inst.eecs.berkeley.edu/webacct`

Once you login using your CalNet ID, you can click on 'Get a new account' in the eecs151 row. Once the account has been created, you can email your class account form to yourself to have a record of your account information.

Logging into the Classroom Servers

The servers used for this class are `c125m-1.eecs.berkeley.edu` through `c125m-16.eecs.berkeley.edu`, and are physically located in Cory 125. You can access all of these machines remotely through SSH. Others such as `eda-1.eecs.berkeley.edu` through `eda-8.eecs.berkeley.edu` are also available for remote login.

To begin this lab, get the project files by typing the following command

```
git clone /home/ff/eecs151/labs/lab1
cd lab1
```

Using Text Editors

Most of the time you will spend designing chips will be spent writing scripts in a text editor. Therefore becoming proficient at editing text is a vital skill. Unlike Java or C programming, there is no integrated development environment (IDE) for writing these scripts. However, many of the advantages of IDE's can be obtained by using the proper editor. In this class, we will be using either Vim or Emacs. Editors such as gedit or nano are not allowed.

If you have never used Vim, please follow the tutorial here: `http://www.openvim.com/tutorial.html` (If you would prefer to learn Emacs, you can read `http://www.gnu.org/software/emacs/tour/` and run the Emacs built-in tutorial with Ctrl-h followed by t) Feel free to search for other resources online to learn more.

Question 1: Common editor tasks

For each task below, describes the keys you need to press to accomplish the following tasks in the file `force_regs.ucli`.

a) Delete 5 lines

b) Search for the text "clock"

c) Replace the text "dut" with "device_under_test"

d) Jump to the end of the file

e) Go to line 42

f) Reload the file (in case it was modified in another window)

g) Save and exit

Linux Basics

You will need to learn how to use linux so that you can understand what programs are running on the server, manipulate files, launch programs, and debug problems. Please read through the tutorial here: `http://linuxcommand.org/lc3_learning_the_shell.php`

To use the CAD tools in this class, you will need to load the class environment. All of the tools are already installed on the network filesystem, but by default users do not have the tools in their path.

Try locating a program that is already installed (`vim`) and another which is not (`icc_shell`).

```
which vim
which unknown_command
```

The program has been installed in: `/usr/bin/vim`. If you show the contents of `/usr/bin`, you will notice that you can launch any of programs by typing their filename. This is because `/usr/bin` is in the environment variable `$PATH`, which contains different directories to search in a colon-separated list.

```
echo $PATH
```

To be able to access the CAD tools, you will need to append to their location to this variable:

```
echo $PATH
source /home/ff/eecs151/tutorials/eecs151.bashrc
which icc_shell
```

Question 2: Common terminal tasks

For each task below, submit the command needed to generate the desired result.

a) List the 5 most recently modified items in `/usr/bin`

b) What directory is `git` installed in?

c) Show the hidden files in your lab directory (the one you cloned from `/home/ff/eecs151/labs/lab1`

d) What version of Vim is installed? Describe how you figured this out.

e) Run `ping www.google.com`, suspend it, then kill the process. Then run it in the background, report its PID, then kill the process.

f) Run `top` and report the average CPU load, the highest CPU job, and the amount of memory used

g) Copy the files in this lab to `/scratch`

There are a few miscellaneous commands to analyze disk usage on the servers.

```
du -ch --max-depth=1 .
df -H
```

Also, it might be helpful to add some features to Bash (the terminal you are using). Try adding the following lines to your .bashrc and restart your session. Now when you change directories, you no longer need to type "ls" to show the directory contents.

```
function cd {
    builtin cd "$@" && ls -F
}
```

Regular Expressions

Regular expressions allow you to perform complex 'Search' or 'Search and Replace' operations. Please work through the tutorial here: `http://regexone.com`

Regular expressions can be used from many different programs: Vim, Emacs, grep, sed, Python, etc. From the command line, use `grep` to search, and `sed` to search and replace.

Unfortunately, deciding what characters needs to be escaped can be somewhat confusing. For example, to find all instances of dcdc_unit_cell_x. where x is a **single digit number**, using grep:

```
grep "unit_cell_[0-9]\{1\}\." force_regs.ucli
```

And you can do the same search in Vim:

```
vim force_regs.ucli
/unit_cell_[0-9]\{1\}\.
```

Notice how you need to be careful what characters get escaped (the [is not escaped but { is). Now imagine we want to add a leading 0 to all of the single digit numbers. The match string in sed could be:

```
sed -e 's/\(unit_cell_\)\([0-9]\{1\}\.\)/\10\2/' force_regs.ucli
```

Both sed, vim, and grep use "Basic Regular Expressions" by default. For regular expressions heavy with special characters, sometimes it makes more sense to assume most characters except a-zA-Z0-9 have special meanings (and they get escaped with only to match them literally). This is called "Extended Regular Expressions", and ?+{}() no longer need to be escaped. A great resource for learning more is `http://en.wikipedia.org/wiki/Regular_expression#POSIX_basic_and_extended`. In Vim, you can do this with \v:

```
:%s/\v(unit_cell_)([0-9]{1}\.)/\10\2/
```

And in sed, you can use the -r flag:

```
sed -r -e 's/(unit_cell_)([0-9]{1}\.)/\10\2/' force_regs.ucli
```

And in grep, you can use the -E flag:

```
grep -E "unit_cell_[0-9]{1}\." force_regs.ucli
```

`sed` and `grep` can be used for many purposes beyond text search and replace. For example, to find all files in the current directory with filenames that contain a specific text string:

```
find . | grep ".ucli"
```

Or to delete all lines in a file that contain a string:

```
sed -e '/reset/d' force_regs.ucli
```

Question 3: Fun with regular expressions

For each regular expression, provide an answer for both basic and extended mode (`sed` and `sed -r`. You are allowed to use multiple commands to perform each task. Operate on the `force_regs.ucli` file.

a) Change all x's surrounding numbers to angle brackets. For example, `regx15xx79x` becomes `reg<15><79>`. Hint: remember to enable global substitution.

b) Make every number in the file be exactly 3 digits with padded leading zeros (except the last 0 on each line). Eg. line 119/120 should read:

```
force -deposit rocketTestHarness.dut.Raven003Top_withoutPads.TileWrap.
...     .io_tilelink_release_data.sync_w002r.rq002_wptr_regx000x.Q 0
force -deposit rocketTestHarness.dut.Raven003Top_withoutPads.TileWrap.
...     .io_tilelink_release_data.fifomem.mem_regx015xx098x.Q 0
```

File Permissions

A tutorial about file permissions can be found here: `http://www.tutorialspoint.com/unix/unix-file-permission.htm`

Question 4: Understaning file permissions

For each task below please provide a command (or multiple) that results in the correct permissions being set. Operate on the `run_always.sh` script.

a) Change the script to be executable by only you

b) Add permission for everyone in your group to be able to execute the same script

c) Make the script writable by you and everyone in your group, but unreadable by others.

d) Change the owner of the file to be eecs151 (Note: You will not be able to run this command, just provide the syntax for it)

Using Makefiles

Makefiles are a simple way to string together a bunch of different shell tasks in an intelligent manner. This allows someone to automate tasks and easily save time when doing repetitive tasks since make targets allow for only files that have changed to need to be updated. Please read through the following tutorial here: `http://www.cs.colby.edu/maxwell/courses/tutorials/maketutor/` (Optional) Further documentation on make can be found here: `http://www.gnu.org/software/make/manual/make.html`

Let's look at a simple makefile to explain a few things about how they work - this is not meant to be anything more than a very brief overview of what a makefile is and how it works. If you look at the Makefile in the provided folder in your favorite text editor, you can see the following lines:

```
output_name = force_regs.random.ucli

$(output_name): force_regs.ucli
    awk 'BEGIN{srand();}{if ($$1 != "") { print $$1,$$2,$$3,int(rand()*2)}}' $< > $@

clean:
    rm -f $(output_name)
```

While this may look like a lot of random characters, let us walk through each part of it to see that it really is not that complicated.

In a Makefile, different chunks of code are run as what are called "targets." The two targets in the above Makefile are `clean` and `output_name`. Here, `output_name` is the name of a variable within the Makefile, which means that it can be overwritten from the command line. This can be done with the following command:

```
make output_name=foo.txt
```

which will result in the output being written to `foo.txt` intstead of `force_regs.random.ucli`.

Inside the `output_name` target, the awk expression has a bunch of $ characters. This is because in normal awk the variable names are `$1, $2`, and then in the makefile you have to escape those variable names to get them to work properly. In make the character to do that is $.

The other characters after the awk script are also special characters to make. The `$<` is the first dependency of that target, the `>` simply redirects the output of awk, and the `$@` is the name of the target itself. This allows users to create makefiles that can be reusable, since you are operating on a dependency and outputting the result into the name of your own target.

The target will run everytime that its dependencies have been updated more recently than its own outputs, so by editing/updating the `force_regs.ucli` file you can get the makefile to rerun, and if there are no edits then it won't change anything. This is different than a bash script, as you can see in `runalways.sh`, which will always run no matter if the file is updated or not.

Question 5: Makefile targets

a) Add a new make target that will create a file called `foo.txt` when that target is run, and will also run the `output_name` target

b) Name at least two ways that you could have the makefile rerun the `output_name` target after it has been run

---

Comparing Files

One indispensable debugging technique is comparing text files. The tools generally behave as black boxes, so during debugging you will be comparing output files to prior output files, and relating that to changes in your input files.

From the command lines, you can use `diff` to compare files:

```
diff force_regs.ucli force_regs.random.ucli
```

You can also compare the contents of directories (the `-q` flag will summarize the results to only show the names of the files that differ, and the `-r` flag will recurse through subdirectories).

For Vim users, there is a useful built-in diff tool:

```
vimdiff force_regs.ucli force_regs.random.ucli
```

Version Control with Git

Version control systems help track how files change overtime and make it easier for collaborators to work on the same files and share their changes. We use `git` to distribute the lab files so that bug fixes can easily be incorporated into your files. Please go through the following tutorial: `try.github.com`

Question 6: Checking Git Understanding

Submit the command required to perform the following tasks.

a) What is the difference between your current Makefile and the file you started with?

b) How do you make a new branch?

c) What is the SHA of the version you checked out?

Customization

Many of the commands and tools you will use on a daily basis can be customized. This can dramatically improve your productivity if used correctly and frequently. Some tools (e.g. `vim` and `bash`) are customized using "dotfiles," which are hidden files in your home directory (e.g. `.bashrc` and `.vimrc`) that contain a series of commands which set variables, create aliases, or change settings. The following links are useful for learning how to make some common customizations. You should read these but are not required to turn in anything for this section.

```
https://www.digitalocean.com/community/tutorials/an-introduction-to-useful-bash-aliases-
and-functions
```

```
http://statico.github.io/vim.html
```

Remote Access

You will find it very useful to be able to continue working on lab (and project!) material from home. There are two convenient ways to remotely access our lab machines: SSH (Secure SHell)

and X2Go.

First, find the name of the workstation you use in the lab. It will look like `c125m-X`, where X is a number from 1 to 23. The *fully qualified DNS name* (FQDN) of your machine is then `c125m-X.eecs.berkeley.edu`. For example, if you were sitting on machine `c125m-8`, the FQDN would be `c125m-8.eecs.berkeley.edu`. Note that you can use any lab machine but that our lab machines aren't very powerful; if everyone uses the same one, everyone will find that their jobs perform poorly. ASIC design tools are resource intensive and will not run well when there are too many simultaneous users on these machines. The easiest way to distribute the load is to have you use the machine you physically use in lab.

Next, note your instructional class acccount name - the one that looks like `eecs151-YYY`, for example `eecs151-abc`. This is the account you created at the start of this lab.

SSH

Linux, BSD, macOS

SSH is the de facto remote terminal tool for Linux and BSD systems (which includes macOS). It lets you login to a text console from anywhere (as long as you have network connectivity). SSH also comes as a standard utility in almost all Linux and BSD systems.

If you're in Linux or BSD, you should be able to access your workstation through SSH by running:

```
ssh eecs151-YYY@c125m-X.eecs.berkeley.edu
```

In our examples, this would be:

```
ssh eecs151-abc@c125m-8.eecs.berkeley.edu
```

Windows

The classic way to use SSH on Windows is to download PuTTY (from here: `https://www.putty.org/`). Download it and login with the FQDN above as the Host and your instructional account username.

Advanced users may wish to install SSH in Cygwin or the Windows Subsystem for Linux and use SSH through there.

X2Go

If you prefer a graphical interface (which feels like sitting at the machine in the lab), you can use X2Go. This is a faster alternative to more traditional X-Forwarding, which you could also do, but which we don't recommend.

Download the X2Go client for your platform from the website: `https://wiki.x2go.org/doku.php/download:start`

To use X2Go, you need to create a new session (look under the Session menu). Give the session any name, it doesn't matter, but set the Host field to the FQDN of your lab machine and the User field to your instructional account username. For "Session type" , select "GNOME". Here's an example from macOS:

Ackowlegement

This lab is the result of the work of many EECS151/251 GSIs over the years including:

Written By:

- Nathan Narevsky (2014, 2017)
- Brian Zimmer (2014)

Modified By:

- John Wright (2015,2016)
- Ali Moin (2017)
- Arya Reais-Parsi (2019)
- Cem Yalcin (2019)